



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Teaching Professor
Lisa Yan

RISC-V Data Transfer

IS AI-GENERATED ART FAIR?

This year, the Colorado State Fair's annual art competition gave out prizes in all the usual categories: painting, quilting, sculpture. But one entrant, Jason M. Allen of Pueblo West, Colo., didn't make his entry with a brush or a lump of clay. He created it with Midjourney, an artificial intelligence program that turns lines of text into hyper-realistic graphics. Mr. Allen's work, "Théâtre D'opéra Spatial," took home the blue ribbon in the fair's contest for emerging digital artists — making it one of the first A.I.-generated pieces to win such a prize, and setting off a fierce backlash from artists who accused him of, essentially, cheating. Reached by phone on Wednesday, Mr. Allen defended his work. He said that he had made clear that his work — which was submitted under the name "Jason M. Allen via Midjourney" — was created using A.I., and that he hadn't deceived anyone about its origins. "I'm not going to apologize for it," he said. "I won, and I didn't break any rules."



Storing Data in Memory

RV32 So Far...

- Addition/subtraction

add rd, rs1, rs2

$R[rd] = R[rs1] + R[rs2]$

sub rd, rs1, rs2

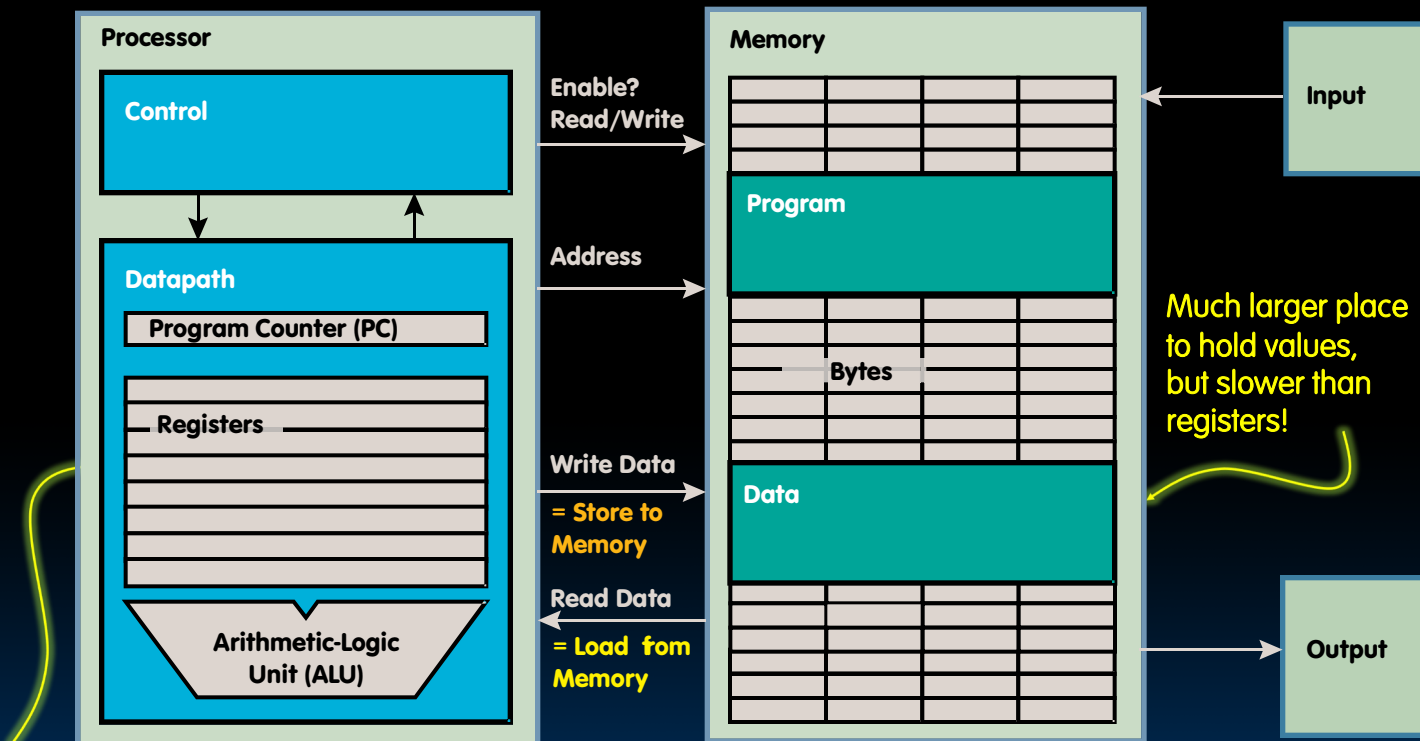
$R[rd] = R[rs1] - R[rs2]$

- Add immediate

addi rd, rs1, imm

$R[rd] = R[rs1] + imm$

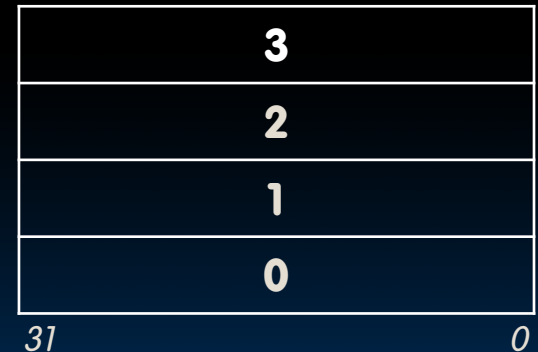
Data Transfer: **Load from** and **Store to** memory



Very fast,
but limited space to hold values!

Memory Addresses are in Bytes

- Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)—works fine if everything is a multiple of 8 bits
- 8 bit chunk is called a *byte* (1 word = 4 bytes)
- Memory addresses are really in *bytes*, not words
- Word addresses are 4 bytes apart
 - Word address is same as address of rightmost byte – least-significant byte (i.e. **Little-endian** convention)



Memory Addresses are in Bytes

- Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)—works fine if everything is a multiple of 8 bits
- 8 bit chunk is called a *byte* (1 word = 4 bytes)
- Memory addresses are really in *bytes*, not words
- Word addresses are 4 bytes apart
 - Word address is same as address of rightmost byte – least-significant byte (i.e. **Little-endian** convention)

Least-significant byte
in a word ↓

15	14	13	12
11	10	9	8
7	6	5	4
3	2	1	0

31 24 23 16 15 8 7 0

Least-significant byte
gets the smallest address



Big Endian vs. Little Endian

The adjective endian has its origin in the writings of 18th century writer Jonathan Swift. In the 1726 novel Gulliver's Travels, he portrays the conflict between sects of Lilliputians divided into those breaking the shell of a boiled egg from the big end or from the little end. He called them the "Big-Endians" and the "Little-Endians".

- The order in which BYTES are stored in memory
- Bits always stored as usual within a byte (E.g., $0xC2 = 0b\ 1100\ 0010$)

Consider the number 1025 as we typically write it:

BYTE3	BYTE2	BYTE1	BYTE0
00000000	00000000	00000100	00000001

Big Endian

ADDR3	ADDR2	ADDR1	ADDR0
BYTE0	BYTE1	BYTE2	BYTE3
00000001	00000100	00000000	00000000

Examples

Names in China or Hungary (e.g., Garcia Dan)

Java Packages: (e.g., org.mypackage.HelloWorld)

Dates in ISO 8601 YYYY-MM-DD (e.g., 2020-09-07)

Eating Pizza crust first

Little Endian

ADDR3	ADDR2	ADDR1	ADDR0
BYTE3	BYTE2	BYTE1	BYTE0
00000000	00000000	00000100	00000001

Examples

Names in the US (e.g., Dan Garcia)

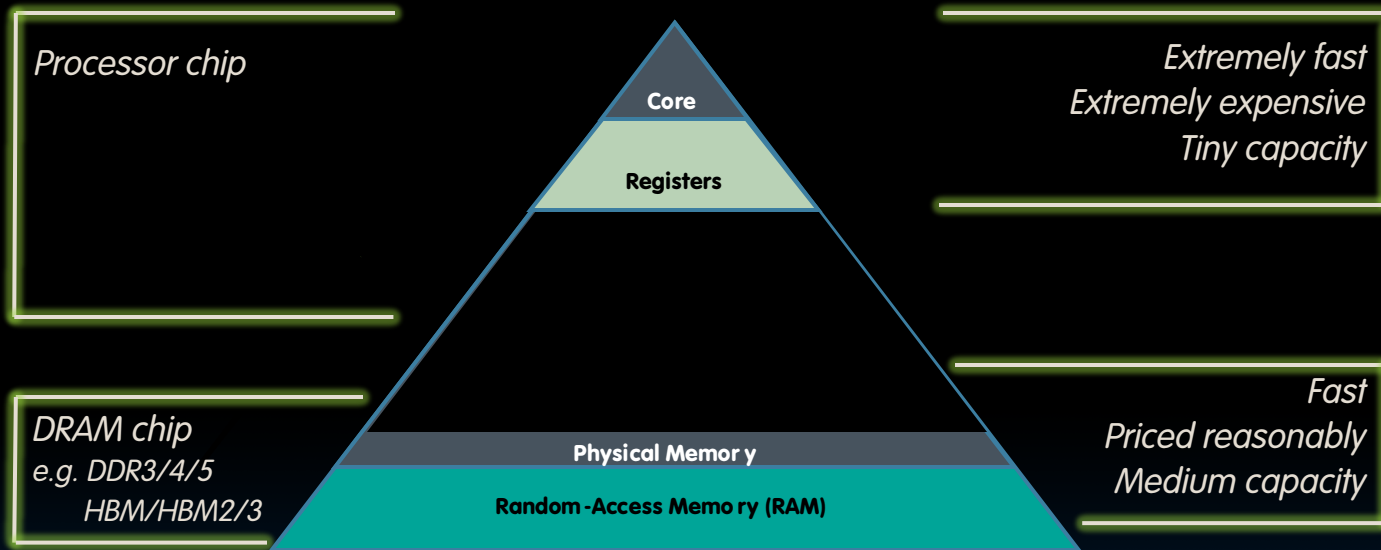
Internet names (e.g., cs.berkeley.edu)

Dates written in Europe DD/MM/YYYY (e.g., 07/09/2020)

Eating Pizza skinny part first

Data Transfer Instructions

Great Idea #3: Principle of Locality / Memory Hierarchy

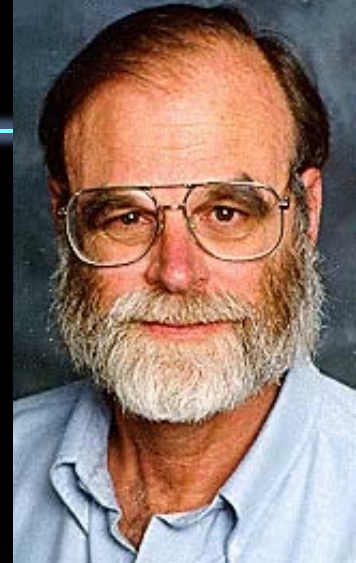




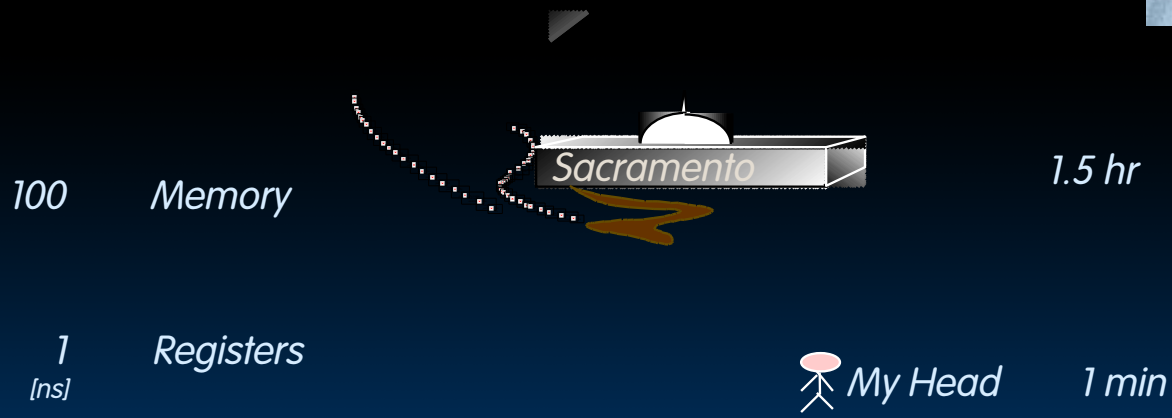
Speed of Registers vs. Memory

- *Given that*
 - *Registers: 32 words (128 Bytes)*
 - *Memory (DRAM): Billions of bytes (2 GB to 64 GB on laptop)*
- *and physics dictates...*
 - *Smaller is faster*
- *How much faster are registers than DRAM??*
 - *About 50-500 times faster! (in terms of **latency** of one access - tens of ns)*
 - *But subsequent words come every few ns*

Jim Gray's Storage Latency Analogy: How Far Away is the Data?



Jim Gray
Turing Award
B.S. Cal 1966
Ph.D. Cal 1969



Load from Memory to Register

- C code

```
int  A[100];
g = h + A[3];
```



- Using Load Word (**lw**) in RISC-V:

```
lw  x10, 12 (x15) # Reg x10 gets A[3]
add x11, x12, x10 # g = h + A[3]
```

Note: **x15** – base register (pointer to A[0])

12 – offset in bytes

Offset must be a constant known at assembly time

Store from Register to Memory

- C code

```
int  A[100];
A[10] = h + A[3];
```

- Using Store Word (**sw**) in RISC-V:

```
lw  x10,12(x15)  # Temp reg x10 gets A[3]
add x10,x12,x10  # Temp reg x10 gets h + A[3]
sw   x10,40(x15) # A[10] = h + A[3]
```



Note: **x15** – base register (pointer)

12,40 – offsets in bytes

x15+12 and **x15+40** must be multiples of 4

Loading and Storing Bytes

- In addition to word data transfers (**lw**, **sw**), RISC-V has **byte** data transfers:

- load byte: **lb**
- store byte: **sb**

- Same format as **lw**, **sw**

- E.g., **lb x10, 3(x11)**

- contents of memory location with address = sum of "3" + contents of register **x11** is copied to the low byte position of register **x10**.

RISC-V also has "unsigned byte" loads (**lbu**) which zero extends to fill register. Why no unsigned store byte '**sbu**'?

x10:

xxxx xxxx xxxx xxxx xxxx **xxxx zzzz zzzz**

...is copied to "sign-extend"

This bit

byte loaded



Example: What is in x12 ?

addi x11,x0,0x93F5

sw x11,0(x5)

lb x12,1(x5)

x5 

x11 
x12 

Memory



L08a What is in x12?

What is in x12 ?

```
addi x11,x0,0x93F5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

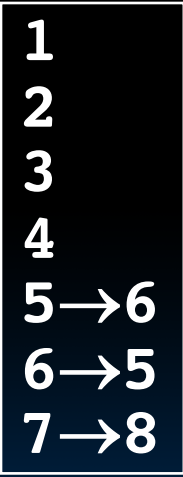
0x0
0x1
0x12
0xF5
0x93
0x9300
0x93F5
0x99999993
0x99999345
0xFFFF9345
0xFFFFF93
0xFFFFFFFF

Example: Translate $*x = *y$

We want to translate $*x = *y$ into RISC-V
 x, y ptrs stored in: $x3$ $x5$

```

1: add x3, x5, zero
2: add x5, x3, zero
3: lw x3, 0(x5)
4: lw x5, 0(x3)
5: lw x8, 0(x5)
6: sw x8, 0(x3)
7: lw x5, 0(x8)
8: sw x3, 0(x8)
  
```



L08b Translate $*x = *y;$

We want to translate $*x = *y$ into RISC-V
 x, y ptrs stored in: $x3$ $x5$

```
1: add x3, x5, zero
2: add x5, x3, zero
3: lw x3, 0(x5)
4: lw x5, 0(x3)
5: lw x8, 0(x5)
6: sw x8, 0(x3)
7: lw x5, 0(x8)
8: sw x3, 0(x8)
```

1
2
3
4
5 —> 6
6 —> 5
7 —> 8



Substituting `addi`

The following two instructions:

```
lw  x10, 12(x15)  # Temp reg x10 gets A[3]
add x12, x12, x10  # reg x12 = reg x12 + A[3]
```

Replace `addi`:

```
addi x12, value # value in A[3]
```

But involve a load from memory!

Add immediate is so common that it deserves its own instruction!

And in Conclusion...

- Memory is **byte**-addressable, but **lw** and **sw** access one **word** at a time.
- A pointer (used by **lw** and **sw**) is just a memory address, we can add to it or subtract from it (using offset).
- **Big- vs Little Endian**
 - **Tip: draw lowest byte on the right**
- **New Instructions:**
lw, sw, lb, sb, lbu