

# Goalkeeping Robot Dog Tends Net Like a Pro



The Mini Cheetah robotic goalie developed by scientists in the [University of California, Berkeley's Hybrid Robotics Lab](#) can save 87.5% of shots taken on goal, compared to almost 80% achieved by players in the English Premier League. The researchers designed the quadrupedal robot to integrate a locomotion controller with a planner for the end-effector trajectory to find the best way to maneuver itself in front of the ball, in less than a second. They trained Mini Cheetah on goalkeeping skills such as sidestepping for intercepts nearby and close to the ground; diving to reach the goal's lower corners, and jumping to cover the top of the goal and the upper corners.

<https://spectrum.ieee.org/football-garcia, Yan-robot-mini-cheetah>





UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley  
Teaching Professor  
Lisa Yan

## Caches II

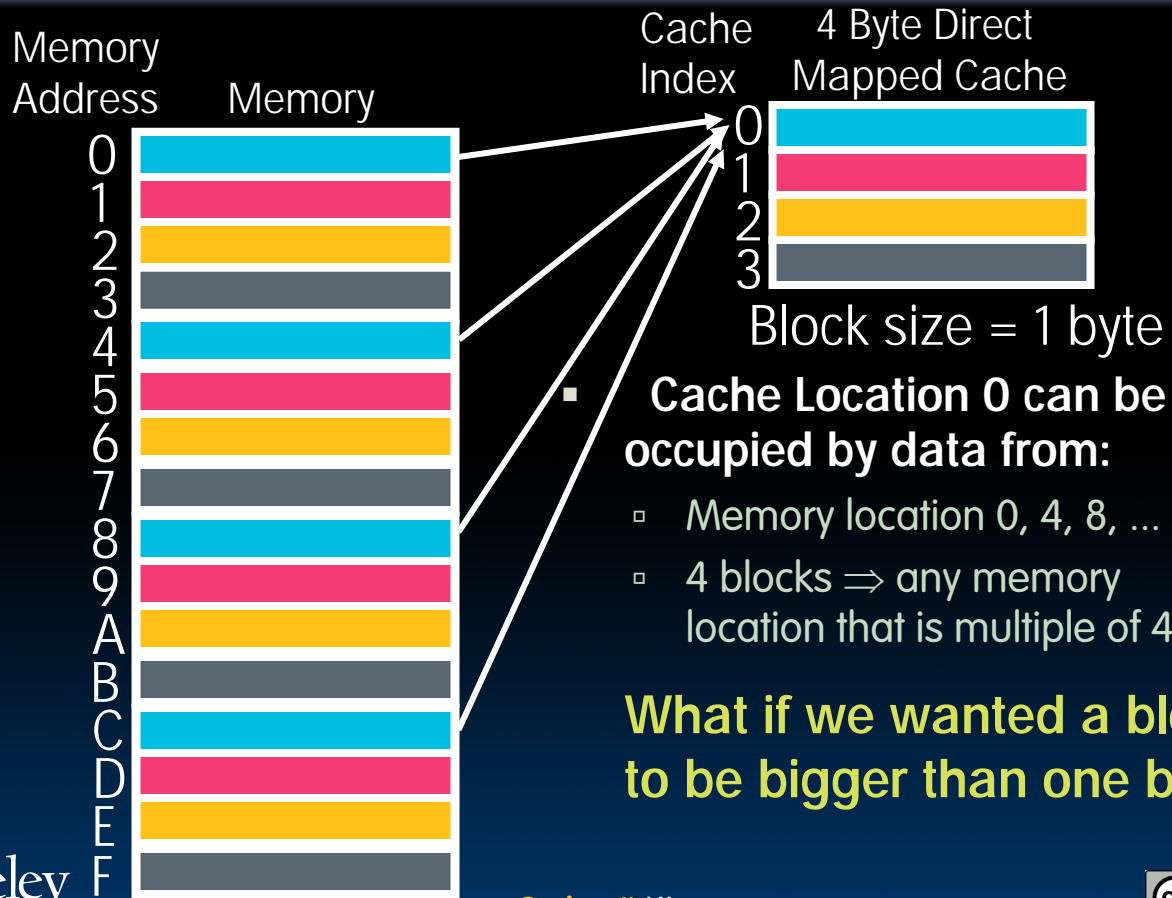
# Direct Mapped Caches

# Direct-Mapped Cache (1/4)

---

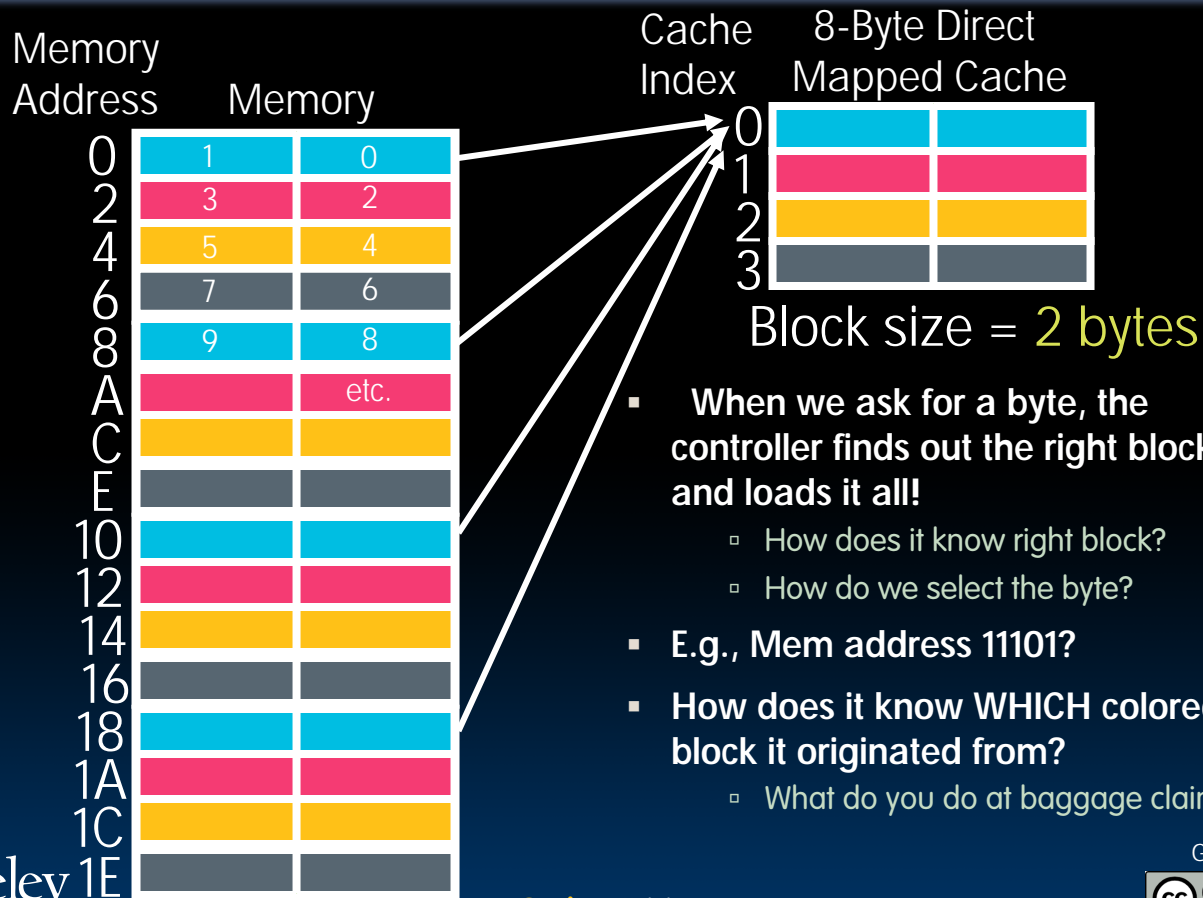
- In a **direct-mapped cache**, each memory address is associated with one possible **block** within the cache
  - Therefore, we only need to look in a single location in the cache for the data if it exists in the cache
  - Block is the unit of transfer between cache and memory

# Direct-Mapped Cache (2/4)

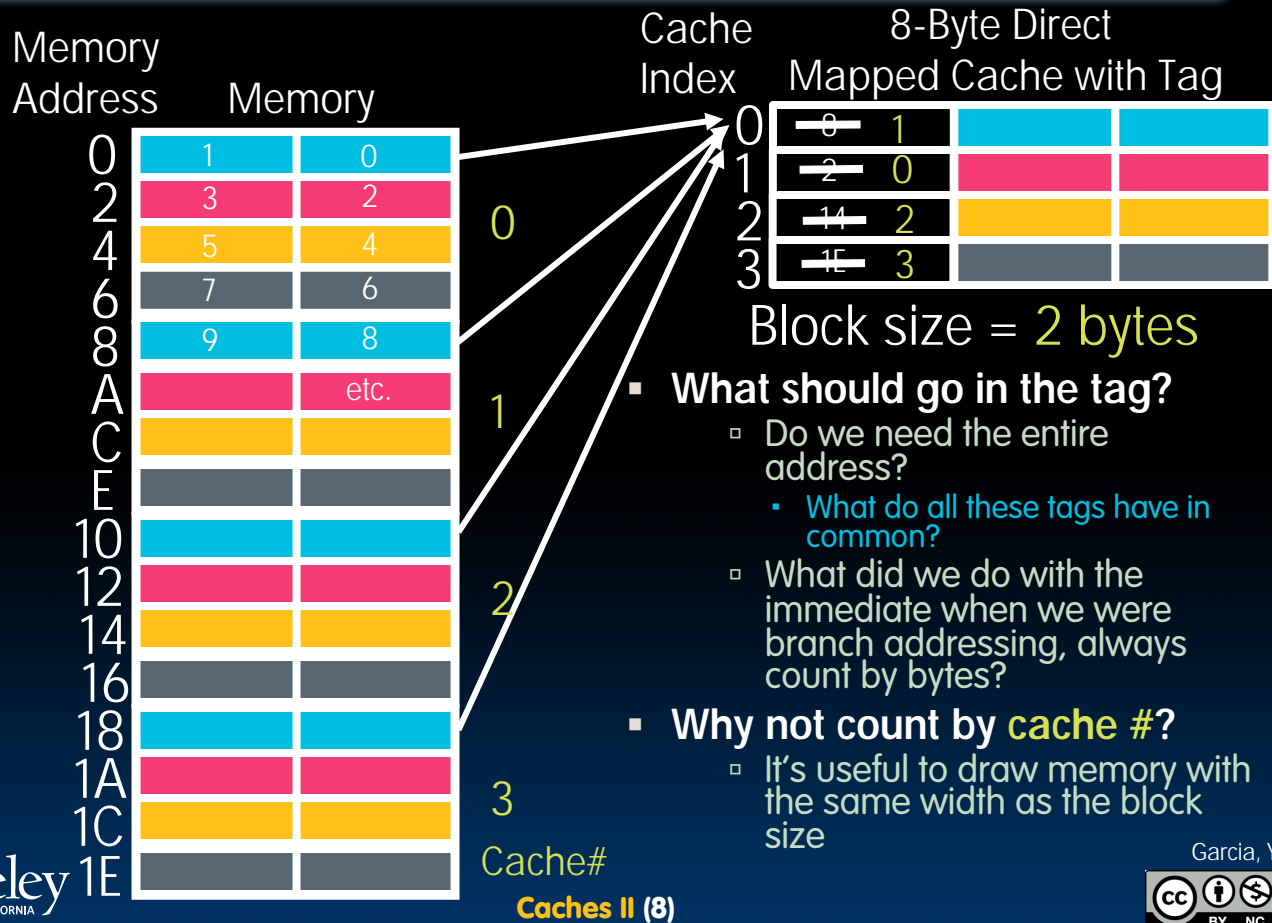


**What if we wanted a block to be bigger than one byte?**

# Direct-Mapped Cache (3/4)



# Direct-Mapped Cache (4/4)





# Issues with Direct-Mapped

- Since multiple memory addresses map to same cache index, how do we tell which one is in there?
- What if we have a block size  $> 1$  byte?
- Answer: divide memory address into three fields





# Direct-Mapped Cache Terminology

---

- All fields are read as unsigned integers.
- **Index**
  - specifies the cache index (which “row”/block of the cache we should look in)
- **Offset**
  - once we’ve found correct block, specifies which byte within the block we want
- **Tag**
  - the remaining bits after offset and index are determined; these are used to distinguish between all the memory addresses that map to the same location



# IQ Cache Mnemonic (Thanks Uncle Dan!)

AREA (cache size, B)

= HEIGHT (# of blocks)


\* WIDTH (size of one block, B/block)

$$2^{(H+W)} = 2^H * 2^W$$

Tag Index Offset

HEIGHT  
(# of blocks)

WIDTH  
(size of one block, B/block)



# Direct Mapped Example



# Direct-Mapped Cache Example (1/3)

- Suppose we have a 8B of data in a direct-mapped cache with 2-byte blocks
  - Sound familiar?
- Determine the size of the tag, index and offset fields if using a 32-bit arch (RV32)
- Offset
  - need to specify correct byte within a block
  - block contains 2 bytes
    - =  $2^1$  bytes
  - need 1 bit to specify correct byte



# Direct-Mapped Cache Example (2/3)

- Index: (~index into an “array of blocks”)
  - need to specify correct block in cache
  - cache contains  $8 \text{ B} = 2^3 \text{ bytes}$
  - block contains  $2 \text{ B} = 2^1 \text{ bytes}$
  - # blocks/cache
    - =  $\frac{\text{bytes/cache}}{\text{bytes/block}}$
    - =  $\frac{2^3 \text{ bytes/cache}}{2^1 \text{ bytes/block}}$
    - =  $2^2 \text{ blocks/cache}$
  - need **2 bits** to specify this many blocks

# Direct-Mapped Cache Example (3/3)

- **Tag: use remaining bits as tag**
  - tag length = addr length – offset - index
    - = 32 - 1 - 2 bits
    - = 29 bits
  - so tag is leftmost **29 bits** of memory address
  - Tag can be thought of as “cache number”
- **Why not full 32-bit address as tag?**
  - All bytes within block need same address
  - Index must be same for every address within a block, so it's redundant in tag check, thus can leave off to save memory

# Memory Access without Cache

- Load word instruction: `lw t0, 0(t1)`
- `t1` contains  $1022_{\text{ten}}$ , `Memory[1022] = 99`
  1. Processor issues address  $1022_{\text{ten}}$  to Memory
  2. Memory reads word at address  $1022_{\text{ten}}$  (99)
  3. Memory sends 99 to Processor
  4. Processor loads 99 into register `t0`



# Memory Access with Cache

- Load word instruction:  $\text{lw } t0, 0(t1)$
- $t1$  contains  $1022_{\text{ten}}$ ,  $\text{Memory}[1022] = 99$
- With cache (similar to a hash)
  1. Processor issues address  $1022_{\text{ten}}$  to Cache
  2. Cache checks to see if has copy of data at address  $1022_{\text{ten}}$ 
    - 2a. If finds a match (Hit): cache reads 99, sends to processor
    - 2b. No match (Miss): cache sends address 1022 to Memory
      - I. Memory reads 99 at address  $1022_{\text{ten}}$
      - II. Memory sends 99 to Cache
      - III. Cache replaces word with new 99
      - IV. Cache sends 99 to processor
  3. Processor loads 99 into register  $t0$

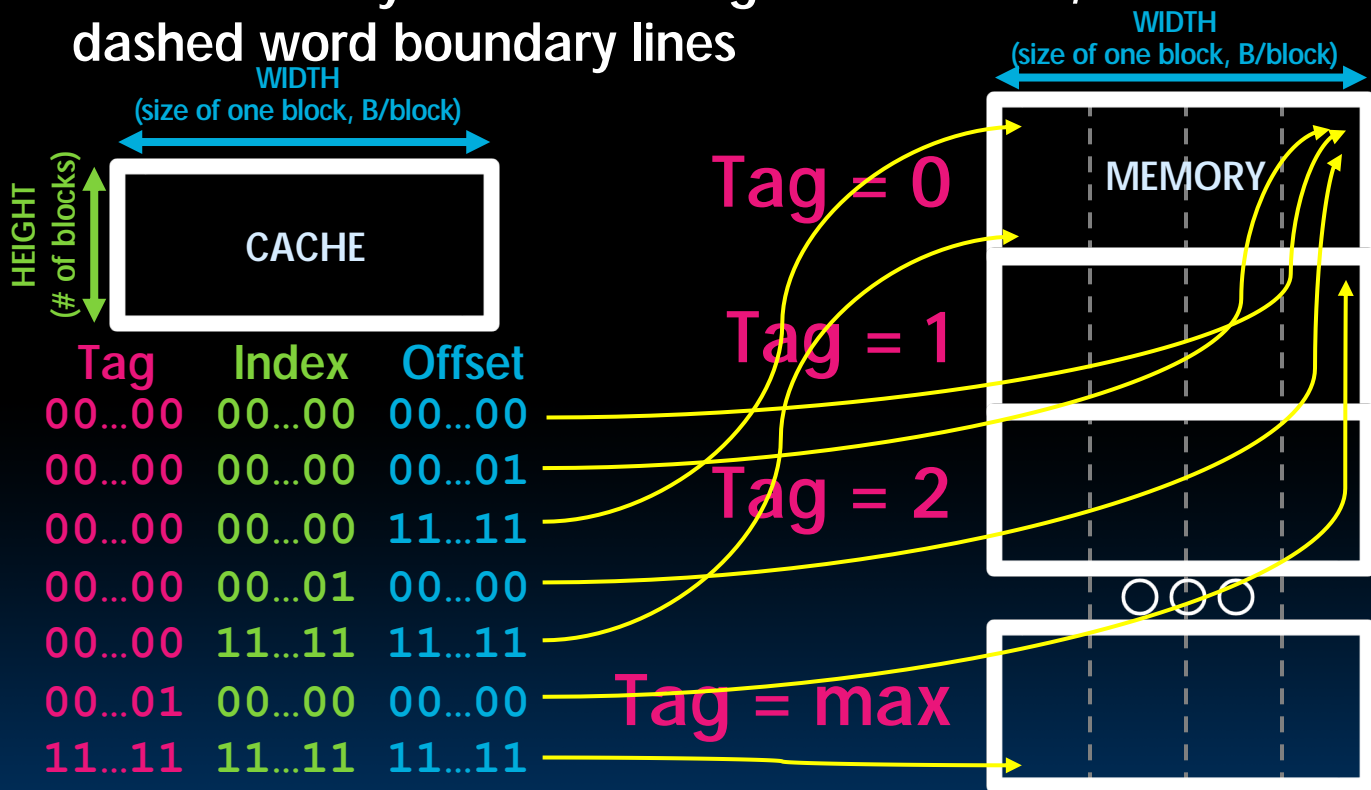
# Solving Cache problems

Tag

Index

Offset

- Draw memory a block wide given T | O bits, dashed word boundary lines



# Cache Terminology

# Caching Terminology

- When reading memory, 3 things can happen:
  - **cache hit:**  
cache block is valid and contains proper address, so read desired word
  - **cache miss:**  
nothing in cache in appropriate block, so fetch from memory
  - **cache miss, block replacement:**  
wrong data is in cache at appropriate block, so discard it and fetch desired data from memory (cache always copy)



# Cache Temperatures

---

- **Cold**
  - Cache empty
- **Warming**
  - Cache filling with values you'll hopefully be accessing again soon
- **Warm**
  - Cache is doing its job, fair % of hits
- **Hot**
  - Cache is doing very well, high % of hits

# Cache Terms

- **Hit rate**: fraction of access that hit in the cache
- **Miss rate**:  $1 - \text{Hit rate}$
- **Miss penalty**: time to replace a block from lower level in memory hierarchy to cache
- **Hit time**: time to access cache memory (including tag comparison)
- **Abbreviation**: "\$" = cache
  - ...a Berkeley innovation!

# One More Detail: Valid Bit

- When start a new program, cache does not have valid information for this program
- Need an indicator whether this tag entry is valid for this program
- Add a “valid bit” to the cache tag entry
  - 0 → cache miss, even if by chance, address = tag
  - 1 → cache hit, if processor address = tag

# Example: 16 KB Direct-Mapped Cache, 16B blocks

- **Valid bit:** determines whether anything is stored in that row (when computer initially powered up, all entries invalid)

**Valid**

Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0				
1	0				
2	0				
3	0				
4	0				
5	0				
6	0				
7	0				
...			...		
1022	0				
1023	0				

Looks like a real cache, will investigate it some more!



L25 For a given cache size: a larger block size can cause a lower hit rate than a smaller one

You can often make your code run faster if you know your computer's cache size

Spatial locality means keep the most recent data closer to the processor

FFF  
FFT  
FTF  
FTT  
TFF  
TFT  
TTF  
TTT

# “And in Conclusion...”

- We have learned the operation of a **direct-mapped cache**
- Mechanism for transparent movement of data among levels of a memory hierarchy
  - set of address/value bindings
  - address → index to set of candidates
  - compare desired address with tag
  - service hit or miss
- load new block and binding on miss

